

# Assignment 2: Ground the Domain - From Naive RAG to Production Patterns

Start Assignment

- Due Friday by 11:59pm
- Points 100
- Submitting a text entry box, a website url, a media recording, or a file upload

## Ground the Domain - From Naive RAG to Production Patterns

### 1. Assignment Overview

**Objective:** Develop practical expertise in Retrieval-Augmented Generation (RAG) systems by implementing naive RAG and two production-ready enhancements, then evaluating performance using industry-standard metrics to understand real-world RAG deployment challenges.

**Core Learning Goals:** By completing this assignment, you will demonstrate proficiency in RAG pipeline architecture, advanced retrieval techniques, systematic evaluation using RAGAs/ARES frameworks, and evidence-based optimization of retrieval-augmented systems for production deployment.

### 2. Resource Access and Equity

#### Supported Platforms (Free/Accessible)

- **Large Language Models:** - **Claude.ai** (free tier) - **ChatGPT-3.5** (free tier) - **Google Gemini** (free tier) - **Llama 3.1/3.2** via Ollama (local installation) - **Mistral 7B** via Ollama or Hugging Face - **Qwen 2.5** via Ollama or Hugging Face Transformers - **Flan-T5** (various sizes) via Hugging Face Transformers
- **Embedding Models:** - **sentence-transformers** (all-MiniLM-L6-v2, all-mpnet-base-v2) - **Hugging Face sentence embeddings** (free tier) - **text-embedding-ada-002** (if using free OpenAI credits)
- **Vector Databases:** - **FAISS** (Facebook AI Similarity Search - completely free) - **Milvus Lite** (local, lightweight version) - **Chroma** (open-source vector database) - **Weaviate** (free tier for small datasets)
- **Computational Platforms:** - **Google Colab** (free tier with GPU access) - **Kaggle Notebooks** (free GPU hours) - **Local development** (CPU-based implementations supported)

#### Alternative Platforms (Optional)

**If you have personal subscriptions or free credits**, you may use: - **ChatGPT-4o, GPT-4, or GPT-4 Turbo** (OpenAI) - **Claude 3.5 Sonnet or Claude Pro** (Anthropic) - **Gemini Pro or Gemini Advanced** (Google) - **Commercial APIs**: OpenAI, Anthropic, Cohere, or Perplexity Pro - **Cloud Computing**: AWS, Google Cloud, or Azure free/student credits - **Premium Vector Databases**: Pinecone, Qdrant Cloud, or Weaviate Cloud

## Equity Policy

- **No Required Purchases**: Students are **NOT required** to obtain premium versions of LLMs, cloud services, or computational resources
- **No Penalties**: Students will **NOT be penalized** for using open-source, open-weight, or free-tier models instead of premium alternatives
- **Equal Assessment**: All submissions will be evaluated based on methodology, analysis quality, and technical understanding—not on the sophistication of paid tools used

## Dataset Access

- **Primary Dataset**: [RAG Mini Wikipedia](https://huggingface.co/datasets/rag-datasets/rag-mini-wikipedia)  (<https://huggingface.co/datasets/rag-datasets/rag-mini-wikipedia>)

# 3. Progressive Assignment Structure

## Step 1: Domain Documents (Days 1-2)

**Deliverable**: Dataset Setup and Exploration Report (300 words)

### Requirements:

1. **Dataset Access**: Download and explore the [RAG Mini Wikipedia dataset](https://huggingface.co/datasets/rag-datasets/rag-mini-wikipedia)  (<https://huggingface.co/datasets/rag-datasets/rag-mini-wikipedia>)
2. **Data Understanding**: Document dataset structure, sample entries, and data quality observations
3. **Infrastructure Planning**: Set up development environment and verify access to chosen platforms

## Step 2: Build Naive RAG (Days 3-5)

**Deliverable**: Functional RAG Pipeline + Documentation (500 words + code)

### Technical Requirements:

1. **Embed with sentence-transformers**: Use “all-MiniLM-L6-v2” as recommended
2. **Store embeddings in vector DB**: Use Milvus Lite or FAISS as recommended
3. **Search and generate**: Implement retrieval and response generation

**Code Organization**: - Modular implementation with clear documentation - Error handling and logging  
- Configuration files for easy parameter adjustment

## Step 3: Evaluation Phase I (Days 6-8)

**Deliverable:** Initial Evaluation Results (400 words + results)

**Requirements:**

1. **Generate answers:** Pass **top-1** retrieved document to prompt
2. **Prompting strategies:** Use different approaches (CoT, Persona Prompting, Instruction Prompt, etc.)
3. **Calculate metrics:** F1-score and Exact Match using HuggingFace Squad metric 4.
4. **Document best strategy:** Note which prompting strategy performed best with hypothesis

## Step 4: Experimentation (Days 9-11)

**Deliverable:** Parameter Comparison Analysis (500 words + results table)

**Experimental Requirements:**

1. **Embedding sizes:** Generate results using **at least 2 different embedding sizes** (e.g., 256, 384, 512)
2. **Retrieval variations:** Experiment with top 3/5/10 retrieved documents (**at least 2 different passage selection strategies** beyond top-1)
3. **Repeat evaluation:** Apply Step 3 evaluation methodology to all parameter combinations
4. **Documentation:** Submit results as markdown/PDF document in GitHub repository

## Step 5: Add Two Advanced RAG Features (Days 12-13)

**Deliverable:** Enhanced RAG Implementation (400 words + code)

**Enhancement Options** (from Merwane Merabet list): - Query rewriting - Reranking  
- Metadata filtering - Multi-vector retrieval - Confidence scoring - Context window optimization - Grounded citations

**Implementation Requirements:** - Choose any two enhancements from the list above - Implement both features and integrate with existing pipeline - Rerun test queries on enhanced system

## Step 6: Advanced Evaluation with RAGAs or ARES (Days 14-15)

**Deliverable:** Automated Evaluation Report (500 words + metrics)

**Evaluation Requirements:** - Use either RAGAs or ARES to assess naive and enhanced RAG pipelines - Report on **at least** the following metrics: - Faithfulness - Context precision/recall - Answer helpfulness or relevance - Compare performance between naive and enhanced systems

## Step 7: Synthesis and Reflection (Days 16-17)

**Deliverable:** Final Summary Report (1000-1250 words)

**Report Structure:** - Description of naive system with all required details - Experimentation results from Steps 3 and 4 - Rationale for chosen enhancements  
- Results from RAGAs/ARES evaluation - Key lessons, limitations, and potential for further improvement

**Additional Requirements:** - Professional formatting appropriate for technical audience - Evidence-based conclusions with supporting data - AI usage documentation as per course policy

## 4. Enhanced Academic Integrity and AI Collaboration Framework

### Comprehensive Usage Documentation

**Required Logging** (consistent with Assignment 1): - **Tool Inventory:** All GenAI tools, models, APIs, and libraries used - **Usage Log:** Specific purposes, input queries, and output utilization - **Code Generation:** Any AI-assisted code with clear attribution - **Content Creation:** AI involvement in documentation, analysis, or reporting

### Permitted AI Assistance

**Acceptable Uses:** - Code debugging and optimization suggestions - Technical documentation assistance - Literature review and reference formatting - Performance analysis and interpretation guidance

#### Required Attribution Format:

##### ## AI Usage Log

```
- **Tool**:[Name and Version]
- **Purpose**:[Specific task]
- **Input**:[Query or prompt provided]
- **Output Usage**:[How AI response was incorporated]
- **Verification**:[How accuracy was confirmed]
```

### Prohibited Practices

- Using AI to implement core RAG components without understanding
- Submitting AI-generated analysis as original insights
- Using AI to interpret evaluation results without independent verification
- Undocumented collaboration with AI systems

### Academic Integrity Verification

**Required Evidence:** - Personal contribution statements for each major component - Decision rationale documentation showing independent reasoning - Technical troubleshooting logs demonstrating hands-on problem solving - Original analysis connecting results to broader RAG concepts

## 5. Detailed Technical Specifications

### Reproducibility Requirements

#### Code Organization:

```
assignment2-rag/  
├── src/  
│   ├── naive_rag.py  
│   ├── enhanced_rag.py  
│   ├── evaluation.py  
│   └── utils.py  
├── data/  
│   ├── processed/  
│   └── evaluation/  
├── results/  
│   ├── naive_results.json  
│   ├── enhanced_results.json  
│   └── comparison_analysis.csv  
├── notebooks/  
│   ├── data_exploration.ipynb  
│   ├── system_evaluation.ipynb  
│   └── final_analysis.ipynb  
├── docs/  
│   ├── setup_instructions.md  
│   ├── evaluation_report.pdf  
│   └── technical_appendix.md  
└── requirements.txt
```

**Documentation Standards:** - Comprehensive README with setup instructions - Inline code comments explaining complex logic - Configuration files with parameter explanations - Error handling and logging throughout

### Evaluation Methodology Standards

**Minimum Testing Requirements:** - **Query Diversity:** Minimum 100 test queries across different complexity levels - **Statistical Significance:** Report confidence intervals where applicable - **Failure Analysis:** Systematic categorization of error types - **Baseline Comparison:** Clear improvement measurements over naive RAG

## 6. Deliverables and Assessment Structure

### Required Submissions

Component	Format	Word Count/Size	Weight	Due Phase
Architecture Document	.md or .pdf	750 words	10%	Phase 1
Naive RAG Implementation	.py/.ipynb + report	500 words + code	25%	Phase 2
Enhanced RAG System	.py/.ipynb + analysis	750 words + code	25%	Phase 3
Evaluation Report	.pdf	1000 words	20%	Phase 4
Technical Report	.pdf	1250-1500 words	15%	Phase 5
Complete GitHub Repository	URL	N/A	5%	Final
<b>Total</b>			<b>100%</b>	

### Technical Report Structure (Phase 5)

**Executive Summary** (200 words): Key findings, performance improvements, deployment readiness

**System Architecture** (300 words): Technical implementation details, design decisions, trade-offs

**Experimental Results** (400-500 words): Comprehensive performance analysis with statistical support

**Enhancement Analysis** (300-400 words): Advanced technique effectiveness, implementation challenges

**Production Considerations** (200-300 words): Scalability, deployment recommendations, limitations

**Appendices**: Complete AI usage logs, technical specifications, reproducibility instructions

## 7. Aligned Evaluation Rubric

### Excellent (90-100%)

- **Technical Implementation:** Robust, well-documented code with comprehensive error handling
- **Experimental Design:** Systematic approach with appropriate controls and statistical analysis
- **Enhancement Integration:** Sophisticated implementation with measurable performance improvements
- **Evaluation Rigor:** Thorough analysis using multiple metrics with insightful interpretation
- **Professional Communication:** Clear, comprehensive documentation appropriate for technical audience

### Proficient (80-89%)

- **Technical Implementation:** Functional code with adequate documentation and basic error handling
- **Experimental Design:** Sound methodology with appropriate comparisons and basic statistical analysis
- **Enhancement Integration:** Competent implementation with demonstrable improvements
- **Evaluation Rigor:** Adequate analysis using required metrics with reasonable interpretation
- **Professional Communication:** Generally clear documentation with minor gaps

### Developing (70-79%)

- **Technical Implementation:** Basic functionality with limited documentation or error handling
- **Experimental Design:** Minimal methodology with superficial comparisons
- **Enhancement Integration:** Incomplete or poorly integrated enhancements with unclear impact
- **Evaluation Rigor:** Partial analysis with limited metric usage or interpretation
- **Professional Communication:** Unclear documentation with significant gaps

### Needs Improvement (Below 70%)

- **Technical Implementation:** Non-functional or poorly implemented code
- **Experimental Design:** Inadequate methodology without proper controls
- **Enhancement Integration:** Missing or ineffective enhancements
- **Evaluation Rigor:** Minimal or absent evaluation with no meaningful analysis
- **Professional Communication:** Poor documentation hindering reproducibility

## 8. Timeline and Resources

### Recommended Timeline

**Estimated Total Time:** 18-24 hours over 17 days

**Days 1-2:** Dataset setup and exploration (2 hours)

**Days 3-5:** Naive RAG implementation (6 hours)

**Days 6-8:** Initial evaluation and prompting experiments (3 hours)

**Days 9-11:** Parameter experimentation across embedding sizes and retrieval counts (4 hours)

**Days 12-13:** Advanced RAG enhancements implementation (5 hours)

**Days 14-15:** RAGAs/ARES evaluation (3 hours)

**Days 16-17:** Final synthesis report and documentation (3 hours)

## Technical Resources

**Required Tools and Libraries:** - Python 3.8+ with virtual environment - Core libraries: sentence-transformers, FAISS/Milvus, transformers - RAGAs or ARES evaluation frameworks - Jupyter notebooks for experimentation and analysis - Git for version control

**Recommended Development Setup:** - Local Python environment with GPU support (if available) - Google Colab or Kaggle for cloud-based development - VS Code or PyCharm for code development - GitHub for repository hosting and submission

## 9. Enhanced Tool and Resource Specifications




### Recommended Technology Stack

**Core Libraries** (with fallback options): - **Vector Storage:** Milvus Lite (primary), FAISS (fallback) - **Embeddings:** sentence-transformers (primary), OpenAI embeddings (if available) - **LLMs:** Llama 2/3, Qwen, Flan-T5 (primary), OpenAI GPT-3.5/4 (if available) - **Evaluation:** RAGAs (primary), ARES (alternative), custom metrics

**Development Environment:** - Python 3.8+ with virtual environment - Jupyter notebooks for experimentation and analysis - Git for version control with meaningful commits - Requirements.txt with pinned versions for reproducibility

### Updated Reading List

**Required Reading** (building on Assignment 1):

1. Merabet, M. "[17 Advanced RAG Techniques to Turn Your RAG App Prototype into a Production-Ready Solution](https://towardsdatascience.com/17-advanced-rag-techniques-to-turn-your-rag-app-prototype-into-a-production-ready-solution-5a048e36cdc8/)  [\(https://towardsdatascience.com/17-advanced-rag-techniques-to-turn-your-rag-app-prototype-into-a-production-ready-solution-5a048e36cdc8/\)](https://towardsdatascience.com/17-advanced-rag-techniques-to-turn-your-rag-app-prototype-into-a-production-ready-solution-5a048e36cdc8/)." *Towards Data Science*, 2024.
2. [RAGAs GitHub Repository and Documentation](https://github.com/explodinggradients/ragas)  [\(https://github.com/explodinggradients/ragas\)](https://github.com/explodinggradients/ragas)
3. [ARES: An Automated Evaluation Framework for RAG](https://github.com/stanford-futuredata/ARES)  [\(https://github.com/stanford-futuredata/ARES\)](https://github.com/stanford-futuredata/ARES)



4. [Evaluating RAG Applications with RAGAs](https://medium.com/data-science/evaluating-rag-applications-with-ragas-81d67b0ee31a)  [\(https://medium.com/data-science/evaluating-rag-applications-with-ragas-81d67b0ee31a\)](https://medium.com/data-science/evaluating-rag-applications-with-ragas-81d67b0ee31a) (Medium Article)
5. Alammar & Grootendorst, [Hands-On LLMs](https://learning.oreilly.com/library/view/hands-on-large-language/9781098150952/)  [\(https://learning.oreilly.com/library/view/hands-on-large-language/9781098150952/\)](https://learning.oreilly.com/library/view/hands-on-large-language/9781098150952/) (Ch. 6–7)
6. Bouchard & Peters, [Building LLMs for Production](https://learning.oreilly.com/library/view/building-llms-for/9798324731472/)  [\(https://learning.oreilly.com/library/view/building-llms-for/9798324731472/\)](https://learning.oreilly.com/library/view/building-llms-for/9798324731472/) (Ch. 4)

**Supplementary Resources:** - Vector database comparison guides - Embedding model benchmarks and selection criteria - Production RAG deployment case studies - Evaluation metric interpretation guides

## 10. Connection to Course Progression

### Building on Assignment 1

- **Domain Expertise:** Apply domain knowledge from Assignment 1 to RAG system design
- **Evaluation Skills:** Extend prompt evaluation experience to RAG-specific metrics
- **Systematic Analysis:** Build on comparative methodology skills


### Preparing for Assignment 3

- **Technical Foundation:** RAG pipeline serves as base for reasoning enhancement
- **Evaluation Framework:** Metrics and methodology extend to fine-tuning assessment
- **Production Readiness:** Advanced techniques prepare for deployment considerations

### Portfolio Integration

- Document technical growth and learning progression
- Build comprehensive case study of LLM application development
- Develop professional portfolio materials for industry applications

**Link to the Starter Code for this assignment -**

<https://github.com/sayakbanerjee1999/Application-of-LLM-Assignment-2-Starter-Code>   
(<https://github.com/sayakbanerjee1999/Application-of-LLM-Assignment-2-Starter-Code>)